

Stage 3 Part 2

— Query 1 —

```
EXPLAIN ANALYZE
SELECT p.Plan_ID, p.Plan_Name, u.User_ID, u.Name, pc.Total_Sessions
FROM Plan p
JOIN Users u ON p.User_ID = u.User_ID
JOIN (
    SELECT Plan_ID, COUNT(Session_ID) AS Total_Sessions
    FROM Plan_Contains
    GROUP BY Plan_ID
) pc ON p.Plan_ID = pc.Plan_ID
ORDER BY pc.Total_Sessions DESC
LIMIT 15;
```

Explain Analyze Output

```
| -> Limit: 15 row(s) (cost=13463.00 rows=15) (actual time=8.24
    -> Nested loop inner join (cost=13463.00 rows=5700) (actual
        -> Nested loop inner join (cost=11468.00 rows=5700) (actual
            -> Sort: pc.Total_Sessions DESC (cost=9473.00..9473.00
                -> Table scan on pc (cost=1717.51..1791.25 rows=5700)
                    -> Materialize (cost=1717.50..1717.50 rows=5700)
                        -> Group aggregate: count(Plan_Contains
                            -> Covering index scan on Plan_Contains
                                -> Filter: (p.User_ID is not null) (cost=0.25 rows=15)
                                    -> Single-row index lookup on p using PRIMARY (Plan_ID)
                                        -> Single-row index lookup on u using PRIMARY (User_ID=
|
```

Analysis

Costs: 13463.00

Execution Time: 8.280ms..10.74ms

Total

- **High Cost in Sorting and Grouping:**

- The GROUP BY in the subquery where it counts sessions per plan, is costing a lot per the output above
- Sorting also contributed a lot to the execution time by processing with 0 indexing

- **Joins on User_ID and Plan_ID:**

- We have Nested loop inner joins which is a indicator that indexing User_ID and Plan_ID in Plan and Plan_Contains respectively, should improve performance by avoiding full scans

Post-Analysis Additions

We are going to index these two columns.

```
CREATE INDEX idx_plan_user_id ON Plan(User_ID);  
CREATE INDEX idx_plan_contains_plan_id ON Plan_Contains(Plan_ID);
```

```
mysql> CREATE INDEX idx_plan_user_id ON Plan(User_ID);  
Query OK, 0 rows affected (0.46 sec)  
Records: 0 Duplicates: 0 Warnings: 0  
  
mysql> CREATE INDEX idx_plan_contains_plan_id ON Plan_Contains(Plan_ID);  
Query OK, 0 rows affected (0.24 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

New Explain Analyze Output

```
| -> Limit: 15 row(s) (cost=13463.00 rows=15) (actual time=4.30  
      -> Nested loop inner join (cost=13463.00 rows=5700) (actual  
            -> Nested loop inner join (cost=11468.00 rows=5700) (actual  
                  -> Sort: pc.Total_Sessions DESC (cost=9473.00..9473.00)
```

```

-> Table scan on pc (cost=1717.51..1791.25 rows=
  -> Materialize (cost=1717.50..1717.50 rows=
    -> Group aggregate: count(Plan_Contains
      -> Covering index scan on Plan_Contai
    -> Filter: (p.User_ID is not null) (cost=0.25 rows=
      -> Single-row index lookup on p using PRIMARY (f
    -> Single-row index lookup on u using PRIMARY (User_ID=|

```

Post Analysis

- **Execution Time :**

- **Before Indexing:** 8.280..10.741 ms
- **After Indexing:** 4.307..5.552 ms

The added indexes significantly reduced query time.

- **Optimized Grouping and Sorting:**

- With idx_plan_contains_plan_id, the Group aggregate in the subquery now uses an indexed scan instead of a table scan.
- **Cost Improvement:** Group aggregate on Plan_Contains.Session_ID improved from 0.308..2.258 ms to 0.046..1.326 ms.
- This addition now allows faster sorting and grouping operations in Plan_Contains.

- **Improved Join Efficiency:**

- The new indexes allows for faster lookups for User_ID in Plan and Plan_ID in Plan_Contains, this allows MySQL to use indexed lookups instead of table scans, resulting in a more efficient query.

— Query 2 —

```

EXPLAIN ANALYZE
WITH RankedExercises AS (
    SELECT e.Exercise_Name, e.Muscle_Group, e.Difficulty,
           ROW_NUMBER() OVER (PARTITION BY e.Muscle_Group ORDER
FROM Exercises e
JOIN Sets s ON e.Exercise_ID = s.Set_ID
JOIN Users u ON s.User_ID = u.User_ID
WHERE e.Difficulty > (
    SELECT AVG(e2.Difficulty)
    FROM Exercises e2
    WHERE e2.Muscle_Group = e.Muscle_Group
)
)
SELECT Exercise_Name, Muscle_Group, Difficulty
FROM RankedExercises
WHERE rn = 1
ORDER BY Difficulty DESC
LIMIT 15;

```

OUTPUT

```

| -> Limit: 15 row(s) (cost=0.10..0.10 rows=0) (actual time=23:
    -> Sort: RankedExercises.Difficulty DESC, limit input to 15
        -> Index lookup on RankedExercises using <auto_key0> (rn
            -> Materialize CTE RankedExercises (cost=0.00..0.00
                -> Window aggregate: row_number() OVER (PARTITI
                    -> Sort: e.Muscle_Group, e.Difficulty DESC
                        -> Stream results (cost=2337.10 rows=29
                            -> Nested loop inner join (cost=23:
                                -> Nested loop inner join (cost
                                    -> Filter: (e.Difficulty >
                                        -> Table scan on e (cos
                                            -> Select #3 (subquery :
                                                -> Aggregate: avg(e:

```

```

-> Filter: (e2.I
      -> Table scan
-> Filter: (s.User_ID is not null)
      -> Single-row index lookup
-> Single-row covering index lookup

```

Analysis

Costs: 2337.10

Execution Time: 2331.449 ms .. 2331.450 ms

- Window Aggregate and Sort Operations:
 - The row_number() window function has a high cost, along with all of the sorting operations of muscle group and difficulty
- Subquery:
 - The subquery that calculates the average difficulty for each Muscle_Group also has a high cost.
 - The subquery is executed for each row in the main query, resulting in a large number of rows and loops (2917 loops).
 - A table scan on Exercises (e2) occurs in the subquery, which further increased the cost.

Post Analysis Addition

```

CREATE INDEX idx_muscle_difficulty ON Exercises (Muscle_Group, Difficulty);
CREATE INDEX idx_muscle_group ON Exercises (Muscle_Group);

```

New Explain Analyze Output

```

| -> Limit: 15 row(s) (cost=0.10..0.10 rows=0) (actual time=38.100..38.100 rows=15)
      -> Sort: RankedExercises.Difficulty DESC, limit input to 15
      -> Index lookup on RankedExercises using <auto_key0> (rows=15)

```

```

-> Materialize CTE RankedExercises (cost=0.00..0.00)
    -> Window aggregate: row_number() OVER (PARTITION BY e.Muscle_Group, e.Difficulty ORDER BY e.Muscle_Group, e.Difficulty DESC)
    -> Sort: e.Muscle_Group, e.Difficulty DESC
    -> Stream results (cost=2337.10 rows=2337)
        -> Nested loop inner join (cost=2337.10 rows=2337)
            -> Nested loop inner join (cost=2337.10 rows=2337)
                -> Filter: (e.Difficulty > s.Difficulty)
                    -> Table scan on e (cost=0.00 rows=1000)
                    -> Select #3 (subquery cost=26.06 rows=172)
                        -> Aggregate: avg(e.Difficulty)
                            -> Covering index: idx_muscle_difficulty
                        -> Filter: (s.User_ID is not null)
                            -> Single-row index look up: idx_muscle_group_difficulty
                        -> Single-row covering index look up: idx_muscle_group_difficulty

```

Post Analysis

- **Reduced Cost in the Window Aggregate and Sort Operations**

- The addition of idx_muscle_difficulty and idx_muscle_group helps in efficiently partitioning and sorting:
 - **Window aggregate:** 2,300 ms to 385.601..386.261 ms.
 - **Sort on e.Muscle_Group and e.Difficulty:** 385.590..385.732 ms.
- The index on (Muscle_Group, Difficulty) optimized these operations by allowing MySQL to efficiently sort and partition the data based on indexed values.

- **Improved Efficiency in Select #3**

- The subquery now uses a covering index (idx_muscle_difficulty) to look up Muscle_Group and Difficulty, reducing the subquery cost:
 - Now **cost=26.06 and rows=172** with actual time= 0.017..0.104 ms per loop, significantly lower than before.

— Query 3 —

```
EXPLAIN ANALYZE
SELECT p.Plan_ID, p.Plan_Name, COUNT(DISTINCT e.Muscle_Group) AS Muscle_Groups_Targeted
FROM Plan p
JOIN Plan_Contains pc ON p.Plan_ID = pc.Plan_ID
JOIN Session s ON pc.Session_ID = s.Session_ID
JOIN Session_Contains sc ON s.Session_ID = sc.Session_ID
JOIN Sets se ON sc.Set_ID = se.Set_ID
JOIN Set_Contains sec ON se.Set_ID = sec.Set_ID
JOIN Exercises e ON sec.Exercise_ID = e.Exercise_ID
GROUP BY p.Plan_ID, p.Plan_Name
ORDER BY Muscle_Groups_Targeted DESC, p.Plan_ID DESC
LIMIT 15;
```

OUTPUT

```
| -> Limit: 15 row(s) (actual time=298.304..298.306 rows=15 loops=1)
      -> Sort: Muscle_Groups_Targeted DESC, p.Plan_ID DESC, limit 15 (cost=0.00..0.00 rows=15 width=16)
            -> Stream results (cost=49824.74 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                  -> Group aggregate: count(distinct e.Muscle_Group) (cost=0.00..0.00 rows=1 width=16)
                        -> Nested loop inner join (cost=44522.42 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                                -> Nested loop inner join (cost=25964.28 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                                        -> Nested loop inner join (cost=16347.00 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                                                -> Nested loop inner join (cost=10000.00 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                                                        -> Nested loop inner join (cost=10000.00 rows=53023) (actual time=298.304..298.306 rows=15 loops=1)
                                                                -> Sort: p.Plan_ID, p.Plan_Name (cost=0.00..0.00 rows=53023 width=16)
                                                                        -> Table scan on p (cost=0.00..0.00 rows=53023 width=16)
                                                                                -> Covering index lookup on pc (cost=0.00..0.00 rows=53023 width=16)
                                                                                        -> Single-row covering index lookup on s (cost=0.00..0.00 rows=53023 width=16)
                                                                                                -> Covering index lookup on sc (cost=0.00..0.00 rows=53023 width=16)
                                                                                                        -> Single-row covering index lookup on se (cost=0.00..0.00 rows=53023 width=16)
                                                                                                                -> Single-row covering index lookup on sec (cost=0.00..0.00 rows=53023 width=16)
                                                                                                                        -> Covering index lookup on sec using Plan_Contains_Plan_ID_idx (cost=0.00..0.00 rows=53023 width=16)
```

-> Single-row index lookup on e using PRIMARY

|

Analysis

- **Total Cost:** 49824.74
- **Execution Time:** 298.304 ms
- The query performs multiple nested loop joins across, each contributing to high costs.
- The GROUP BY and COUNT(DISTINCT e.Muscle_Group) on Exercises would benefit from a sort and aggregate operation

To improve efficiency, let's index columns frequently used in joins

- **Plan_Contains.Plan_ID** (Plan)
- **Session_Contains.Session_ID** (Session and Plan_Contains)
- **Set_Contains.Set_ID** (Sets and Session_Contains)

Post Analysis Addition

```
CREATE INDEX idx_plan_contains_plan_id ON Plan_Contains(Plan_ID);
CREATE INDEX idx_session_contains_session_id ON Session_Contains(Session_ID);
CREATE INDEX idx_set_contains_set_id ON Set_Contains(Set_ID);
```

New Explain Analyze Output

```
| -> Limit: 15 row(s) (actual time=211.554..211.556 rows=15 loops=1)
      -> Sort: Muscle_Groups_Targeted DESC, p.Plan_ID DESC, limit 15
            -> Stream results (cost=46641.21 rows=53023) (actual time=211.554..211.556 rows=15 loops=1)
                  -> Group aggregate: count(distinct e.Muscle_Group)
                        -> Nested loop inner join (cost=41338.88 rows=53023) (actual time=211.554..211.556 rows=15 loops=1)
                              -> Nested loop inner join (cost=22780.75 rows=53023) (actual time=211.554..211.556 rows=15 loops=1)
                                    -> Nested loop inner join (cost=13163.88 rows=53023) (actual time=211.554..211.556 rows=15 loops=1)
                                          -> Nested loop inner join (cost=6841.21 rows=53023) (actual time=211.554..211.556 rows=15 loops=1)
```



```

-> Nested loop inner join (cost
    -> Nested loop inner join (
        -> Sort: p.Plan_ID, p.P
        -> Table scan on p
        -> Covering index lookup
        -> Single-row covering index
    -> Covering index lookup on sc i
    -> Single-row covering index lookup
    -> Covering index lookup on sec using PI
-> Single-row index lookup on e using PRIMAI

```

Post Analysis

COST: Lowers by about 6.4%,

- **Before Indexing:** 49824.74
- **After Indexing:** 46641.21

TIME: Lowers by about 29%,

- **Before Indexing:** 298.304 ms
- **After Indexing:** 211.554 ms

Summary:

Session_Contains.Session_ID and Set_Contains.Set_ID lookups are now using indexes, reducing time from row-by-row access to direct lookups.

- Session_Contains.Session_ID decreased from 0.018 per lookup to 0.003, and Set_Contains.Set_ID from 0.006 to 0.002.

— Query 4 —

```

EXPLAIN ANALYZE
WITH ExerciseFrequency AS (
    SELECT e.Exercise_ID, e.Exercise_Name, e.Muscle_Group, COUNT(*)
    FROM Exercises e

```

```

JOIN Set_Contains sc ON e.Exercise_ID = sc.Exercise_ID
GROUP BY e.Exercise_ID, e.Exercise_Name, e.Muscle_Group
), RankedExercises AS (
    SELECT Exercise_Name, Muscle_Group,
           ROW_NUMBER() OVER (PARTITION BY Muscle_Group ORDER BY
FROM ExerciseFrequency
)
SELECT Exercise_Name, Muscle_Group
FROM RankedExercises
WHERE rn = 1
ORDER BY Muscle_Group
LIMIT 15;

```

OUTPUT

```

| -> Limit: 15 row(s) (cost=0.10..0.10 rows=0) (actual time=209.003 ms)
    -> Sort: RankedExercises.Muscle_Group, limit input to 15 rows
        -> Index lookup on RankedExercises using <auto_key0> (rows=15)
            -> Materialize CTE RankedExercises (cost=0.00..0.00 rows=0)
                -> Window aggregate: row_number() OVER (PARTITION BY Muscle_Group ORDER BY Exercise_ID)
                    -> Sort: ExerciseFrequency.Muscle_Group, Exercise_ID
                        -> Table scan on ExerciseFrequency (cost=0.00..0.00 rows=0)
                            -> Materialize CTE ExerciseFrequency
                                -> Table scan on <temporary> (cost=0.00..0.00 rows=0)
                                    -> Aggregate using temporary
                                        -> Nested loop inner join
                                            -> Table scan on e
                                                -> Covering index lookup on Set_Contains

```

Analysis

- **Total Cost:** 5645.34
- **Execution Time:** 209.003 ms for the limit operation

- Joins on Set_Contains.Exercise_ID and Exercises.Exercise_ID are costly

We should index the columns that are used in JOINS and SORTING operations

- **Set_Contains.Exercise_ID**: Used to join with Exercises on Exercise_ID.
- **Exercises.Muscle_Group** and **Exercises.Exercise_ID**: Used in grouping and ordering operations in the ExerciseFrequency

Post Analysis Addition

```
CREATE INDEX idx_set_contains_exercise_id ON Set_Contains(Exercise_ID);
CREATE INDEX idx_exercises_muscle_group_exercise_id ON Exercises(Exercise_ID, Muscle_Group);
```

New Explain Analyze Output

```
| -> Limit: 15 row(s) (cost=0.10..0.10 rows=0) (actual time=10.000..10.000 rows=15)
    -> Sort: RankedExercises.Muscle_Group, limit input to 15 rows
        -> Index lookup on RankedExercises using <auto_key0> (rows=15)
            -> Materialize CTE RankedExercises (cost=0.00..0.00 rows=0)
                -> Window aggregate: row_number() OVER (PARTITION BY Exercise_ID ORDER BY Exercise_ID)
                    -> Sort: ExerciseFrequency.Muscle_Group, Exercise_ID
                        -> Table scan on ExerciseFrequency (cost=0.00..0.00 rows=0)
                            -> Materialize CTE ExerciseFrequency
                                -> Table scan on <temporary> (cost=0.00..0.00 rows=0)
                                    -> Aggregate using temporary
                                        -> Nested loop inner join
                                            -> Table scan on e
                                                -> Covering index lookup
```

Post Analysis

Total Cost: Cost went up in indexing, I believe this is due to the increased operations that are involved in cost calculations, meaning indexing hurt cost more than it helped.

- **Before Indexing:** 5645.34

- **After Indexing:** 7845.56

Execution Time: Execution time has decreased by approximately 51%, showing greater retrieval speeds to the indexing

- **Before Indexing:** 209.003 ms
- **After Indexing:** 102.212 ms

Nested Joins and Table Access:

- Set_Contains.Exercise_ID lookup decreased to 0.005..0.007 ms from 0.015..0.019 ms .

Summary: Although cost time went up, our execution time went down dramatically